

Safety mean: harmful-content refusal +
jailbreak robustness (higher is better)

*strong on both
(safe & helpful)*

0.975
0.950
0.925
0.900
0.875
0.850
0.825
0.800

0.70

0.75

0.80

0.85

0.90

Over-refusal domain mean
(compliance on benign prompts — higher is better)

K2-Horizon 75

K2-Horizon 36B

K2-Horizon

K2-Horizon 4B

K2-Horizon 1B

